

# SQL

## Spring 2018

School of Computer Science  
University of Waterloo

Databases CS348

# SQL (Structured Query Language)



- Based on the **Relational Calculus**:
  - ⇒ Conjunctive queries  
aka. `SELECT` blocks
  - ⇒ Set operations
  - ⇒ Update language
  - ⇒ Non First-order features
- BAG (multiset) Semantics
- NULL values
  - ⇒ avoid if at all possible
- A **committee** design
  - ⇒ often more **pragmatic** than logical
  - ⇒ evolving *standard*:  
SQL-89, **SQL-92**, SQL-1999,  
SQL:2003/2006/2008/2011/2016

# SQL (cont.)

Three major parts of the language:

**1** DML (Data Manipulation Language)

⇒ Query language

⇒ Update language

**Also:** Embedded SQL (SQL/J) and ODBC (JDBC)

⇒ necessary for application development

**2** DDL (Data Definition Language)

⇒ defines *schema* for relations

⇒ creates (modifies/destroys) database objects.

**3** DCL (Data Control Language)

⇒ access control

# SQL Data Types

## Values of attributes in SQL:

|                            |                                       |
|----------------------------|---------------------------------------|
| <code>integer</code>       | integer (32 bit)                      |
| <code>smallint</code>      | integer (16 bit)                      |
| <code>decimal(m, n)</code> | fixed decimal                         |
| <code>float</code>         | IEEE float (32 bit)                   |
| <code>char(n)</code>       | character string (length $n$ )        |
| <code>varchar(n)</code>    | variable length string (at most $n$ ) |
| <code>date</code>          | year/month/day                        |
| <code>time</code>          | hh:mm:ss.ss                           |

# Sample Database Revisited

```
AUTHOR(aid integer, name char(20))
```

```
WROTE(author integer, publication char(8))
```

```
PUBLICATION(pubid char(8), title char(70))
```

```
BOOK(pubid char(8),  
      publisher char(50), year integer)
```

```
JOURNAL(pubid char(8),  
         volume integer, no integer, year integer)
```

```
PROCEEDINGS(pubid char(8),  
             year integer)
```

```
ARTICLE(pubid char(8), crossref char(8),  
         startpage integer, endpage integer)
```

... SQL is **NOT** case sensitive.

# The Basic “SELECT Block”

Basic syntax:

```
3 SELECT DISTINCT <results>
1 FROM             <tables>
2 WHERE            <condition>
```

- Allows formulation of conjunctive ( $\exists, \wedge$ ) queries of the form

$$\{ \langle \text{results} \rangle \mid \exists \langle \text{unused} \rangle. \left( \bigwedge \langle \text{tables} \rangle \right) \wedge \langle \text{condition} \rangle \}$$

$\Rightarrow$  a conjunction of  $\langle \text{tables} \rangle$  with  $\langle \text{condition} \rangle$

$\Rightarrow \langle \text{results} \rangle$  specifies values in the resulting tuples, and

$\Rightarrow \langle \text{unused} \rangle$  are variables *not used in*  $\langle \text{results} \rangle$

# Example

List all authors in the database:

```
SQL> select distinct *  
      2  from author;
```

```
AID  NAME
```

```
---  -----
```

```
1  Toman, David
```

```
2  Chomicki, Jan
```

```
3  Saake, Gunter
```

The FROM clause cannot be used on its own

⇒ the "SELECT \*" notation

⇒ also reveals all attribute names

# Variables vs. Attributes

- Relational Calculus uses *positional* notation, i.e.,  
 $\text{EMP}(x, y, z)$  is true whenever the  $x$ ,  $y$ , and  $z$  components of an answer can be found as a tuple in the instance of  $\text{EMP}$   
 $\Rightarrow$  no need for *attribute names*  
 $\Rightarrow$  inconvenient for relations with high arity
- SQL uses *corelations* (tuple variables) and *attribute names* to assign *default variable names* to components of tuples:  
 $R \text{ [AS] } p$  in SQL stands for  $R(p.a_1, \dots p.a_n)$  in RC  
where  $a_1, \dots, a_k$  are the *attribute names* declared for  $R$ .



## Example

List all publications with at least two authors,

$$\{p \mid \exists a_1, a_2. \text{WROTE}(a_1, p) \wedge \text{WROTE}(a_2, p) \wedge a_1 \neq a_2\} :$$

```
SQL> select distinct r1.publication
      2  from wrote r1, wrote r2
      3  where r1.publication = r2.publication
      4  and r1.author != r2.author;
```

PUBLICATION

-----

ChSa98

ChTo98

ChTo98a

⇒ cannot *share* a variable ( $p$ ) in the two `WROTE` relations  
need for explicit equality “`r1.publication = r2.publication`”

# Example

List titles of all books,

$$\{t \mid \exists p, b, y. \text{PUBLICATION}(p, t) \wedge \text{BOOK}(p, b, y)\} :$$

```
SQL> select distinct title
      2  from publication, book
      3  where publication.pubid = book.pubid;
```

TITLE

-----  
Logics for Databases and Information Systems

⇒ relations can serve as their own *corelations* when *unambiguous*  
publication stands for “publication publication”, i.e.,  
publication(publication.pubid, publication.title)

# The "FROM" Clause (summary)

Syntax:

$$\text{FROM } R_1 [ [\text{AS}] n_1 ], \dots, R_k [ [\text{AS}] n_k ]$$

- $R_i$  are relation (table) names
- $n_i$  are distinct identifiers
- The clause represents a **conjunction**  $R_1 \wedge \dots \wedge R_k$ 
  - $\Rightarrow$  all variables of  $R_i$ 's are *distinct*
  - $\Rightarrow$  we use (co)relation names to resolve ambiguities
- Cannot appear alone
  - $\Rightarrow$  only as a part of the *select block*

# The "SELECT" Clause

Syntax:

SELECT DISTINCT  $e_1$  [ [AS]  $n_1$  ], ...,  $e_k$  [ [AS]  $n_k$  ]

- 1 Eliminate superfluous attributes from answers ( $\exists$ )
- 2 Form **expressions**:  
 $\Rightarrow$  built-in functions applied to values of attributes
- 3 Give names to attributes in the answer

# Standard Expressions

We can **create** values in the answer tuples using **built-in** functions:

- On numeric types:

$+$ ,  $-$ ,  $*$ ,  $/$ , ... (usual arithmetic)

- On strings:

`||` (concatenation), `substr`, ...

- Constants (of appropriate types)

"`SELECT 1`" is a valid query in SQL-92

- UDF (user defined functions)

**Note:** all attribute names **MUST** be “present” in the `FROM` clause.

## Example

For every article list the number of pages:

```
SQL> select distinct pubid, endpage-startpage+1  
      2  from article;
```

| PUBID   | ENDPAGE-STARTPAGE+1 |
|---------|---------------------|
| -----   | -----               |
| ChTo98  | 40                  |
| ChTo98a | 28                  |
| Tom97   | 19                  |

# Naming Attributes in the Results

Results of queries  $\longleftrightarrow$  Tables

What are the names of attributes in the result of a `SELECT` clause?

- A single attribute: inherits the name
- An expression: implementation dependent

We can—and should—**explicitly** name the resulting attributes:

$\Rightarrow$  "`<expr> AS <id>`" where `<id>` is the new name

## Example

For every article list the number of pages, and name the resulting attributes `id`, `numberofpages`:

```
SQL> select distinct pubid as id,  
2         endpage-startpage+1 as numberofpages  
3   from article;
```

| ID      | NUMBEROFPAGES |
|---------|---------------|
| -----   | -----         |
| ChTo98  | 40            |
| ChTo98a | 28            |
| Tom97   | 19            |



# The "WHERE" Clause

Syntax:

```
WHERE <condition>
```

Additional conditions on tuples that qualify for the answer.

- Standard atomic conditions:

- 1 equality: =, != (on all types)

- 2 order: <, <=, >, >=, <> (on numeric and string types)

- Conditions may involve *expressions*

- ⇒ similar conditions as in the `SELECT` clause

## Example(s)

Find all journals printed since 1997:

```
SQL> select * from journal where year >= 1997;
```

| PUBID    | VOLUME | NO | YEAR |
|----------|--------|----|------|
| JLP-3-98 | 35     | 3  | 1998 |

Find all articles with more than 20 pages:

```
SQL> select * from article  
2 where endpage-startpage > 20;
```

| PUBID   | CROSSREF | STARTPAGE | ENDPAGE |
|---------|----------|-----------|---------|
| ChTo98  | ChSa98   | 31        | 70      |
| ChTo98a | JLP-3-98 | 263       | 290     |

# Boolean Connectives

Atomic conditions can be combined using **boolean connectives**:

- AND (conjunction)
- OR (disjunction)
- NOT (negation)

## Example

List all publications with at least two authors:

```
SQL> select distinct r1.publication
  2   from wrote r1, wrote r2
  3   where r1.publication = r2.publication
  4   and not r1.author = r2.author;
```

PUBLICATION

-----

ChSa98

ChTo98

ChTo98a

# Summary

- Simple SELECT block accounts for many queries  
⇒ all in  $\exists, \wedge$  fragment of relational calculus
- Additional features
  - alternative names for relations
  - expressions and naming in the output
  - built-in atomic predicates and boolean connectives
- Well defined semantics (declarative and operational)

# Complex Queries in SQL

- So far we can write only  $\exists, \wedge$  queries
  - $\Rightarrow$  the SELECT BLOCK queries
  - $\Rightarrow$  not sufficient to cover all safe RC queries
- Remaining connectives:
  - 1  $\vee, \neg$ : are expressed using **set operations**
    - $\Rightarrow$  easy to enforce *range-restriction requirements*
  - 2  $\forall$ : rewrite using negation and  $\exists$ 
    - $\Rightarrow$  the same for  $\rightarrow, \leftrightarrow$ , etc.

# Set Operations at Glance

Answers to *Select Blocks* are **relations** (sets of tuples)

⇒ we can apply **set operations** on them

■ Set union:  $Q_1$  UNION  $Q_2$

⇒ the set of tuples in  $Q_1$  or in  $Q_2$ .

⇒ used to express “or”.

■ Set difference:  $Q_1$  EXCEPT  $Q_2$

⇒ the set of tuples in  $Q_1$  but not in  $Q_2$ .

⇒ used to express “and not”.

■ Set intersection:  $Q_1$  INTERSECT  $Q_2$

⇒ the set of tuples in both  $Q_1$  and  $Q_2$ .

⇒ used to express “and” (redundant, rarely used).

$Q_1$  and  $Q_2$  must have **union-compatible** signatures:

⇒ same number and types of attributes

## Example: Union

List all publication ids for books or journals:

```
SQL> (select distinct pubid from book)
      2 union
      3 (select distinct pubid from journal);
```

PUBID

-----

ChSa98

JLP-3-98



## Example: Set Difference

List all publication ids except those for articles:

```
SQL> (select distinct pubid from publication)
  2  except
  3  (select distinct pubid from article);
```

PUBID

-----

ChSa98

DOOD97

JLP-3-98

# What About Nesting of Queries?

We can use *SELECT Blocks* (and other *set operations*) as arguments of *set operations*.

What if we need to use a **set operation** inside of a **SELECT Block**?

- Can use **distributive laws**
  - $\Rightarrow (A \vee B) \wedge C \equiv (A \wedge C) \vee (B \wedge C)$
  - $\Rightarrow$  often **very** cumbersome
- Nest set operation inside a select block.
  - $\Rightarrow$  *common table expressions*

# Naming (Sub-)queries

## Idea:

Queries denote **relations**. We provide a **naming** mechanism that allows us to assign names to (results of) queries.

⇒ can be used later in place of (base) relations.

## ■ Syntax:

```
WITH foo1 [<opt-schema-1>]
    AS ( <query-1-goes-here> ),
    ...
    foon [<opt-schema-n>]
    AS ( <query-n-goes-here> )
<query-that-uses-foo1-...-foon-as-table-names>
```

## Example

List all publication titles for books or journals:

```
SQL> with bookorjournal(pubid) as
2      ( (select distinct pubid from book)
3        union
4          (select distinct pubid from journal)
5        )
6      select distinct title
7      from publication, bookorjournal
8      where publication.pubid = bookorjournal.pubid;
```

TITLE

-----  
Logics for Databases and Information Systems  
Journal of Logic Programming

# The FROM clause revisited

- Using the `WITH` mechanism is sometimes cumbersome:  
⇒ we don't want to name every subexpression

- SQL-92 allows us to **inline** queries in the `FROM` clause:

`FROM ..., ( <query-here> ) <id>, ...`

⇒ `<id>` stands for the result of `<query-here>`.

⇒ unlike for base relations, `<id>` is **mandatory**.

- In “old” SQL (SQL-89) this does NOT work; *views* were the only option...

# Example

List all publication titles for journals or books:

```
SQL> select distinct title
2   from publication,
3       ( (select distinct pubid from journal)
4         union
5         (select distinct pubid from book) ) jb
6   where publication.pubid = jb.pubid;
```

TITLE

-----  
Logics for Databases and Information Systems  
Journal of Logic Programming

# Can't we just use `OR` instead of `UNION`?

- A **common** mistake:

⇒ use of `OR` in the `WHERE` clause instead of the `UNION` operator

- An incorrect solution:

```
select distinct title
from publication, book, journal
where publication.pubid = book.pubid
or publication.pubid = journal.pubid
```

- Often works, but consider where there are no books.

# Summary on First-Order SQL

- SQL introduced so far captures all of RC (relational calculus)
  - ⇒ optionally with duplicate semantics
  - ⇒ powerful (many queries can be expressed)
  - ⇒ efficient (PTIME, LOGSPACE)
- Shortcomings:
  - ⇒ some queries are hard to write (syntactic sugar)
  - ⇒ no “*counting*” (aggregation)
  - ⇒ no “*path in graph*” (recursion)



# WHERE Subqueries

- Additional (complex) search conditions  
⇒ query-based search predicates
- Advantages
  - simplifies writing queries with negation
- Drawbacks
  - complicated semantics  
(especially when duplicates are involved)
  - **very** easy to make mistakes
- **VERY COMMONLY** used to formulate queries

# Overview of WHERE Subqueries

## ■ Presence/absence of a *single* value in a query

`<attr> IN ( <query> )`

`<attr> NOT IN ( <query> )`

## ■ Relationship of a value to some/all values in a query

`<attr> op SOME ( <query> )`

`<attr> op ALL ( <query> )`

## ■ Emptiness/non-emptiness of a query

`EXISTS ( <query> )`

`NOT EXISTS ( <query> )`

In the first two cases `<query>` must be unary.

## Example: “<attr> in (<query>)”

```
SQL> select distinct title
      2   from publication
      3   where pubid in (select pubid from article);
```

TITLE

-----  
Temporal Logic in Information Systems  
Datalog with Integer Periodicity Constraints  
Point-Based Temporal Extension of Temporal SQL

# “Pure” SQL Equivalence

Nesting in the WHERE clause is mere syntactic sugar:

|   |   |
|---|---|
| <pre>select r.b from r where r.a in (   select b   from s )</pre> | <pre>select r.b from r, (   select distinct b   from s ) as s where r.a = s.b</pre> |
|---|---|

All of the remaining constructs can be rewritten in similar fashion.

## Example: “<attr> not in (<query>)”

All author-publication ids for all publications except books and journals:

```
SQL> select *
      2  from wrote
      3  where publication not in (
      4      ( select pubid from book )
      5      union
      6      ( select pubid from journal ) );
```

| AUTHOR | PUBLICAT |
|--------|----------|
| -----  | -----    |
| 1      | ChTo98   |
| 1      | ChTo98a  |
| 1      | Tom97    |
| 2      | ChTo98   |
| 2      | ChTo98a  |

... search conditions may contain complex queries.

“<attr> not in (<query>)” (cont.)

...another formulation:

```
SQL> select *  
2   from wrote  
3   where publication not in (  
4       select pubid from book  
5   ) and publication not in (  
6       select pubid from journal  
7   )
```

...and may be combined using boolean connectives.

**Example:** “<attr> op SOME/ALL (<query>)”

Find the longest articles (a way expressing max):

```
SQL> select distinct pubid
2   from article
3   where endpage-startpage >= all (
4       select endpage-startpage
5       from   article
6   );
```

```
PUBID
-----
ChTo98
```

“<attr> = some (<query>)”  $\equiv$  “<attr> in (<query>)”

“<attr> <> all (<query>)”  $\equiv$  “<attr> not in (<query>)”

# Parametric Subqueries

- So far, *subqueries* were **independent** of the *main* query
  - ⇒ not correlated
  - ⇒ not much fun (good only for simple queries)
- SQL allows **parametric** (correlated) subqueries.

Parametric subqueries have the form “<query>” mentioning

$\langle attr \rangle_1, \langle attr \rangle_2, \dots$

where  $\langle attr \rangle_i$  is an attribute in the main query.

The **truth** of a predicate defined by a subquery is determined for each substitution (tuple) in the main query:

- 1 instantiate all the parameters and
- 2 check for the truth value as before ...



## Example: “EXISTS (<query>)”

Parametric subqueries are most common for “existential” subqueries:

```
SQL> select *
      2  from wrote r
      3  where exists ( select *
      4                    from wrote s
      5                    where r.publication = s.publication
      6                    and r.author <> s.author );
```

| AUTHOR | PUBLICAT |
|--------|----------|
| 1      | ChTo98   |
| 1      | ChTo98a  |
| 2      | ChTo98   |
| 2      | ChTo98a  |
| 2      | ChSa98   |
| 3      | ChSa98   |

## Example: “NOT EXISTS (<query>)”

It is easy to now complement conditions:

```
SQL> select *
      2  from wrote r
      3  where not exists (
      4      select *
      5      from wrote s
      6      where r.publication = s.publication
      7      and r.author <> s.author
      8  );
```

AUTHOR PUBLICAT

-----

1 Tom97

## Example: “<attr> IN (<query>)”

```
SQL> select *  
2   from wrote r  
3   where publication in (  
4       select publication  
5       from wrote s  
6       where r.author <> s.author  
7   );
```

AUTHOR PUBLICAT

```
-----  
1 ChTo98  
1 ChTo98a  
2 ChTo98  
2 ChTo98a  
2 ChSa98  
3 ChSa98
```

# More levels of Nesting

- WHERE subqueries are **just queries**
  - ⇒ we can nest again and again and ...
  - ⇒ every nested subquery can use attributes from the enclosing queries as parameters.
  - ⇒ correct naming is imperative
- Used to formulate very complex **search conditions**
  - ⇒ attributes present **in the subquery only**  
**CANNOT be used** to construct the result(s).

## Example

List all authors who always publish with someone else:

```
SQL> select distinct a1.name
  2  from author a1, author a2
  3  where not exists (
  4      select *
  5      from    publication p, wrote w1
  6      where  p.pubid = w1.publication
  7            and  a1.aid = w1.author
  8            and  a2.aid not in (
  9                select author
 10                from    wrote
 11                where  publication = p.pubid
 12                    and  author <> a1.aid
 13            )
 14 );
```

# Summary

- WHERE subqueries enable easy formulation of queries of the form

“All x in R such that (a part of) x doesn't appear in S”.

- Subqueries only stand for **WHERE conditions**  
⇒ **CANNOT** be used to produce **results**.
- You can use **input parameters**, but these must be *bound* in the main query
- All of these are just a syntactic sugar and can be expressed using queries nested in the FROM clause
  - but it might be quite hard ...
  - and it is easy to make mistakes (be **very** careful)

# How do we Modify a Database?

- Naive approach:

*DBSTART;*

$r_1 := Q_1(DB);$

*...*

$r_k := Q_k(DB);$

*DBCOMMIT;*

- Not an acceptable solution in practice

# Incremental Updates

## Idea

Tables are large but **updates are small**  $\Rightarrow$  **Incremental updates**

- 1 Insertion of a tuples (`INSERT`)
  - $\Rightarrow$  constant tuple
  - $\Rightarrow$  results of queries
- 2 Deletion of tuples (`DELETE`)
  - $\Rightarrow$  based on match of a condition
- 3 Modification of tuples (`UPDATE`)
  - $\Rightarrow$  allows updating “in place”
  - $\Rightarrow$  based on match of a condition



# SQL Insert

- One constant tuple (or a fixed number):

```
INSERT INTO r [ (a1, ..., ak) ]  
VALUES (v1, ..., vk)
```

⇒ adds tuples  $(v_1, \dots, v_k)$  to  $r$ .

⇒ the type of  $(v_1, \dots, v_k)$  must match the schema definition of  $r$ .

- Multiple tuples (generated by a query):

```
INSERT INTO r ( Q )
```

⇒ adds result of  $Q$  to  $r$

## Example: inserton of a tuple

Add a new author:

```
SQL> insert into author
      2          values (4, 'Niwinski, Damian',
      3                  'zls.mimuw.edu.pl/~niwinski');
```

1 row created.

```
SQL> select distinct aid,name,url from author;
```

| AID | NAME             | URL                        |
|-----|------------------|----------------------------|
| 1   | Toman, David     | db.uwaterloo.ca/~david     |
| 2   | Chomicki, Jan    | cs.monmouth.edu/~chomicki  |
| 3   | Saake, Gunter    |                            |
| 4   | Niwinski, Damian | zls.mimuw.edu.pl/~niwinski |

## Example: use of a query

Add a new author (without looking up author id):

```
SQL> insert into author (  
2      select max(aid)+1, 'Snodgrass, Richard T.',  
3          'www.cs.arizona.edu/people/rts'  
4      from author );  
1 row created.
```

```
SQL> select distinct aid, name from author;
```

```
AID  NAME
```

```
-- -----
```

```
1  Toman, David  
2  Chomicki, Jan  
3  Saake, Gunter  
4  Damian Niwinski  
5  Snodgrass, Richard T.
```

# SQL Delete

## ■ Deletion using a condition:

```
DELETE FROM r  
WHERE cond
```

⇒ deletes **all** tuples that match `cond`.

## ■ Deletion using cursors (later)

⇒ available in embedded SQL

⇒ only way to delete one out of two duplicate tuples

## Example

Delete all publications that are not articles or the collections an article appears in:

```
SQL> delete from publication
      2         where pubid not in (
      3             select pubid
      4             from article
      5         ) and pubid not in (
      6             select crossref
      7             from article
      8         );
```

```
0 rows deleted.
```

# SQL Update

- Two components:

- 1 an update statement (`SET`)

- ⇒ an assignment of values to attributes.

- 2 a search condition (`WHERE`)

- Syntax:

```
UPDATE r
SET    <update statement>
WHERE  <condition>
```

# Example

```
SQL> update author
  2  set      url = 'brics.dk/~david'
  3  where   aid in (
  4          select aid
  5          from author
  6          where name like 'Toman%'
  7          );
```

1 row updated.

```
SQL> select * from author;
```

| AID | NAME         | URL               |
|-----|--------------|-------------------|
| 1   | Toman, David | //brics.dk/~david |
| ... |              |                   |

# Support for Transactions

The DBMS guarantees noninterference (serializability) of all data access requests to tables in a database instance

- Transaction starts with first **access** of the database

⇒ until it sees:

- COMMIT: make changes permanent

```
SQL> commit;  
Commit complete.
```

- or ROLLBACK: discard changes

```
SQL> rollback;  
Rollback complete.
```



# Aggregation

Standard and very useful extension of First-Order Queries.

- Aggregate (column) functions are introduced to
  - ⇒ find number of tuples in a relation
  - ⇒ add values of an attribute (over the whole relation)
  - ⇒ find minimal/maximal values of an attribute
- Can apply to *groups* of tuples that with equal values for (selected) attributes
- Can **NOT** be expressed in Relational Calculus

# Aggregation in SQL

The same in SQL syntax:

```
SELECT    x1, ..., xk, agg1, ..., agg1
FROM      Q
GROUP BY  x1, ..., xk
```

Restrictions:

- All attributes in the `SELECT` clause that are **NOT** in the scope of an aggregate function **MUST** appear in the `GROUP BY` clause.
- `aggi` are of the form `count(y)`, `sum(y)`, `min(y)`, `max(y)`, or `avg(y)` where `y` is an attribute of `Q` (usually not in the `GROUP BY` clause).

# Operational Reading

- 1 Partition the input relation to groups with equal values of **grouping** attributes
- 2 On each of these partitions apply the **aggregate function**
- 3 Collect the results and form the answer

## Example (count)

For each publication, count the number of authors:

```
SQL> select publication, count(author) need new name for  
      2   from wrote count(author)  
      3   group by publication;
```

```
PUBLICAT  COUNT(AUTHOR)
```

```
-----  
ChSa98           2  
ChTo98           2  
ChTo98a          2  
Tom97            1
```

## Example (sum)

For each author, count the number of article pages:

```
SQL> select author, sum(endpage-startpage+1) as pgs
2   from wrote, article
3   where publication=pubid
4   group by author;
```

| AUTHOR | PGS |
|--------|-----|
| 1      | 87  |
| 2      | 68  |

... not quite correct: it doesn't list 0 pages for author 3.

# The HAVING clause

- The `WHERE` clause can't impose conditions on values of aggregates  
⇒ `WHERE` clause has to be used **before** `GROUP BY`
- SQL allows a `HAVING` clause instead  
⇒ like `WHERE`, but for aggregate values. . .
- The aggregate functions used in the `HAVING` clause may be different from those in the `SELECT` clause; the grouping, however, is common.

The `HAVING` clause is mere *SYNTACTIC SUGAR*

... and can be replaced by a nested query and a `WHERE` clause.

## Example

List publications with exactly one author:

```
SQL> select publication, count(author)
      2  from wrote
      3  group by publication
      4  having count(author) = 1;
```

```
PUBLICAT  COUNT(AUTHOR)
```

```
-----
```

```
Tom97          1
```

... This query *can* be written without aggregation.

## Example (revisited.)

For every author, count the number of books and articles:

```
SQL> select distinct aid, name, count(publication)
  2  from author, (
  3    ( select distinct author, publication
  4      from wrote, book
  5      where publication = pubid )
  6    union all
  7    ( select distinct author, publication
  8      from wrote, article
  9      where publication = pubid ) ) ba
 10 where aid = author
 11 group by name, aid;
```



# Summary

- SQL covered so far:

- 1 Simple SELECT BLOCK

- 2 Set operations

- 3 Formulation of complex queries, nesting of queries, and views

- 4 Updating Data

- 5 Aggregation

- This covers pretty much all of the useful SQL DML

- ⇒ the Bad and Ugly coming next ...