

SQL: Ordering Results, Duplicate Semantics and NULL Values

Spring 2018

School of Computer Science
University of Waterloo

Databases CS348

Ordering Results

- No particular ordering on the rows of a table can be assumed when queries are written. (This is important!)
- No particular ordering of rows of an intermediate result in the query can be assumed either.
- However, it is possible to order the final result of a query, using the `ORDER BY` clause at the end of the query.

General form:

`ORDER BY e_1 [Dir_1], ..., e_k [Dir_k]`

where Dir_i is either `ASC` or `DESC`; `ASC` is the default.

Example

List all authors in the database in ascending order of their name:

```
SQL> select distinct *  
      2  from author  
      3  order by name;
```

```
AID  NAME
```

```
---  -----  
      2  Chomiccki, Jan  
      3  Saake, Gunter  
      1  Toman, David
```

Again, the `asc` keyword is optional, and is assumed by default. A *descending* order is obtained with the `desc` keyword. Minor sorts, minor minor sorts, etc., can be added.

Multisets and Duplicates

- SQL has always had a **MULTISET/BAG** semantics rather than a **SET** semantics:
 - ⇒ SQL tables are **multisets** of tuples
 - ⇒ originally for efficiency reasons
- What does “*allows duplicates*” mean?

part
bolt
bolt
bolt
nut
nut



part	cnt
bolt	3
nut	2

How does this impact Queries?

Example (Cheap Quantification–Projection)

EMP	
Name	Dept
Bob	CS
Sue	CS
Fred	PMath
Barb	Stats
Jim	Stats

$\{y | \exists x. \text{EMP}(x, y)\}$

Dept
CS
CS
PMath
Stats
Stats



Dept	cnt
CS	2
PMath	1
Stats	2

Range-restricted Queries for Multisets

Definition (Range restricted formulas with DISTINCT)

A formula (condition) φ is *range restricted* when, for φ_i that are also range restricted, φ has the form

$$\begin{array}{ll} R(x_{i_1}, \dots, x_{i_k}), & \\ \varphi_1 \wedge \varphi_2, & \\ \varphi_1 \wedge (x_i = x_j) & (\{x_i, x_j\} \cap FV(\varphi_1) \neq \emptyset), \\ \exists x_i. \varphi_1 & (x_i \in FV(\varphi_1)), \\ \varphi_1 \vee \varphi_2 & (FV(\varphi_1) = FV(\varphi_2)), \\ \varphi_1 \wedge \neg \varphi_2 & (FV(\varphi_2) = FV(\varphi_1)), \text{ or} \\ \text{DISTINCT}(\varphi). & \end{array}$$

Duplicates and Queries

How do we define what an **answer** to a query is now?

Ideas

- 1 A **finite valuation** can appear **k times** ($k > 0$) as a query answer.
- 2 The number of duplicates is a **function** of the numbers of duplicates in formulas.
- 3 $\mathbf{DB}, \theta, k \models \varphi$ reads “finite valuation θ appears **k times** in φ ’s answer”.

Definition (Multiset Semantics for the Relational Calculus)

$\mathbf{DB}, \theta, 0 \models \varphi$	if $\mathbf{DB}, \theta \not\models \varphi$
$\mathbf{DB}, \theta, k \models R(x_1, \dots, x_k)$	if $(\theta(x_1), \dots, \theta(x_k)) \in \mathbf{R}$ k times
$\mathbf{DB}, \theta, m \cdot n \models \varphi \wedge \psi$	if $\mathbf{DB}, \theta, m \models \varphi$ and $\mathbf{DB}, \theta, n \models \psi$
$\mathbf{DB}, \theta, m \models \varphi \wedge (x_i = x_j)$	if $\mathbf{DB}, \theta, m \models \varphi$ and $\theta(x_i) = \theta(x_j)$
$\mathbf{DB}, \theta, \sum_{v \in D} n_v \models \exists x. \varphi$	if $\mathbf{DB}, \theta[x := v], n_v \models \varphi$
$\mathbf{DB}, \theta, m + n \models \varphi \vee \psi$	if $\mathbf{DB}, \theta, m \models \varphi$ and $\mathbf{DB}, \theta, n \models \psi$
$\mathbf{DB}, \theta, m - n \models \varphi \wedge \neg \psi$	if $\mathbf{DB}, \theta, m \models \varphi$, $\mathbf{DB}, \theta, n \models \psi$ and $m > n$

Duplicates and SQL

Allowing duplicates leads to additional *syntax*.

- *A duplicate elimination operator*

- ⇒ “SELECT DISTINCT x” **v.s.** “SELECT x” in SELECT-blocks

- MULTISSET (BAG) operators

- ⇒ equivalents of *set operations*

- ⇒ but with multiset semantics.

Example

```
SQL> select r1.publication
      2  from wrote r1, wrote r2
      3  where r1.publication=r2.publication
      4         and r1.author<>r2.author;
```

PUBLICAT

ChSa98

ChSa98

ChTo98

ChTo98

ChTo98a

ChTo98a

⇒ for publications with n authors we get $O(n^2)$ answers!

Bag Operations

- Bag union: `UNION ALL`

⇒ additive union: bag containing all in Q_1 and Q_2 .

- Bag difference: `EXCEPT ALL`

⇒ subtractive difference (monus):

⇒ a bag all tuples in Q_1 for which
there is no “matching” tuple in Q_2 .

- Bag intersection: `INTERSECT ALL`

⇒ a bag of all tuples taking the
maximal number common to Q_1 and Q_2

Example

```
SQL> ( select author
2      from wrote, book
3      where publication=pubid )
4 union all
5 ( select author
6      from wrote, article
7      where publication=pubid );
```

AUTHOR

```
-----
2
3
1
2
1
2
1
```

Summary

- SQL covered so far:
 - 1 simple SELECT BLOCK
 - 2 set operations
 - 3 duplicates and multiset operations
 - 4 formulation of complex queries, nesting of queries, and views
 - 5 aggregation
- Note that duplicates in subqueries occurring in `where` clauses will not change the results computed by the top-level query, but that this is not true for subqueries in `with` or `from` clauses.

“Pure” SQL Equivalence, Revisited

Recall how nesting in the WHERE clause is syntactic sugar:

SELECT r.b	SELECT r.b
FROM r	FROM r, (
WHERE r.a IN (SELECT DISTINCT b
SELECT b	FROM s
FROM s) AS s
)	WHERE r.a = s.b

Rewriting does not generally hold if `DISTINCT` is removed.

What is a “null” value?

Phone		
Name	Office	Home
Joe	1234	3456
Sue	1235	?

- *Sue doesn't have home phone (value inapplicable)*
- *Sue has home phone, but we don't know her number*
(value unknown)

Value Inapplicable

- Essentially *poor schema design*.
- Better design:

Office Phone	
Name	Office
Joe	1234
Sue	1235

Home Phone	
Name	Home
Joe	3456

- Queries should behave *as if asked* over the above decomposition.
⇒ (relatively) easy to implement

Value Unknown

Idea

Unknown values can be replaced by any domain value (that satisfies integrity constraints).

⇒ many possibilities (possible worlds)

Phone		
Name	Office	Home
Joe	1234	3456
Sue	1235	?

→

Phone		
Name	Office	Home
Joe	1234	3456
Sue	1235	0000

⋮

Phone		
Name	Office	Home
Joe	1234	3456
Sue	1235	9999

Value Unknown and Queries

How do we answer queries?

Idea

Answers true in *all* possible worlds W of an incomplete D .

Certain Answer

$$Q(D) = \bigcap_{W \text{ world of } D} Q(W)$$

\Rightarrow answer common to all possible worlds.

Is this (computationally) feasible?

\Rightarrow NO (NP-hard to *undecidable* except in trivial cases)

SQL's solution: a (crude) approximation

What can we do with NULLs in SQL?

expressions

- general rule: a NULL as a parameter to an operation makes (should make) the result NULL
- $1 + \text{NULL} \rightarrow \text{NULL}$, 'foo' || NULL $\rightarrow \text{NULL}$, etc.

predicates/comparisons

- three-valued logic (crude approximation of “value unknown”)

set operations

- unique *special value* for *duplicates*

aggregate operations

- doesn't “count” (i.e., “value inapplicable”)

Comparisons Revisited

Idea

Comparisons with a NULL value return UNKNOWN

Example

$1 = 1$	TRUE
$1 = \text{NULL}$	UNKNOWN
$1 = 2$	FALSE

Still short of *proper logical* behaviour:

$$x = 0 \vee x \neq 0$$

should be always **true** (no matter what x is, including `NULL`!), but...

UNKNOWN and Boolean Connectives

Idea

Boolean operations have to handle UNKNOWN

\Rightarrow *extended truth tables for Boolean connectives*

\wedge	T	U	F
T	T	U	F
U	U	U	F
F	F	F	F

\vee	T	U	F
T	T	T	T
U	T	U	U
F	T	U	F

\neg	
T	F
U	U
F	T

... for tuples in which x is assigned the NULL value we get:

$$x = 0 \vee x \neq 0 \rightarrow \text{UNKNOWN} \vee \text{UNKNOWN} \rightarrow \text{UNKNOWN}$$

which is not the same as TRUE.

UNKNOWN in WHERE Clauses

How is this used in a WHERE clause?

- Additional syntax IS TRUE, IS FALSE, and IS UNKNOWN
⇒ WHERE <cond> shorthand for WHERE <cond> IS TRUE
- Special comparison IS NULL

List all authors for which we don't know a URL of their home page:

```
SQL> select aid, name
      2  from author
      3  where url IS NULL;
```

```
      AID NAME
-----
      3 Saake, Gunter
```

Counting NULLS

How do NULLs interact with counting (and aggregates in general)?

- `count (URL)` counts only non-NULL URL's

⇒ `count (*)` counts "rows"

```
db2 => select count(*) as RS, count(url) as US
db2 => from author;
```

RS	US
3	2

1 record(s) selected.

Outer Join

Idea

Allow “NULL-padded” answers that “fail to satisfy” a conjunct in a conjunction

- Extension of syntax for the FROM clause

\Rightarrow FROM R <j-type> JOIN S ON C

\Rightarrow the <j-type> is one of FULL, LEFT, RIGHT, or INNER

- Semantics (for $R(x, y)$, $S(y, z)$, and $C = (r.y = s.y)$).

1 $\{(x, y, z) : R(x, y) \wedge S(y, z)\}$

2 $\{(x, y, \text{NULL}) : R(x, y) \wedge \neg(\exists z. S(y, z))\}$ for LEFT and FULL

3 $\{(\text{NULL}, y, z) : S(y, z) \wedge \neg(\exists x. R(x, y))\}$ for RIGHT and FULL

\Rightarrow syntactic sugar for UNION ALL

Example

```
db2 => select aid,publication
db2 => from author left join wrote
db2 =>                on aid=author;
```

AID	PUBLICATION
-----	-------------

-----	-----
-------	-------

1	ChTo98
---	--------

1	ChTo98a
---	---------

1	Tom97
---	-------

2	ChTo98
---	--------

2	ChTo98a
---	---------

2	ChSa98
---	--------

3	ChSa98
---	--------

5	-
---	---

8 record(s) selected.

Counting with OJ

For every author count the number of publications:

```
db2 => select aid, count(publication) as pubs
db2 => from author left join wrote
db2 =>                on aid=author
db2 => group by aid;
```

AID	PUBS
1	3
2	3
3	1
5	0

4 record(s) selected.

Summary

- `NULL`s are necessary evil
 - ⇒ used to account for (small) irregularities in data
 - ⇒ should be used sparingly
- Can be **always** avoided
 - ⇒ however some of the solutions may be inefficient
- You can't escape `NULL`s in practice
 - ⇒ *easy fix* for blunders in schema design
 - ⇒ ... also due to schema evolution, etc.