

# Translating Entity-Relationship to Relational Tables

Spring 2018

School of Computer Science  
University of Waterloo

Databases CS348

# E-R Diagram to Relational Schema

Main ideas:

- Each entity set maps to a new table
- Each attribute maps to a new table column
- Each relationship set maps to either new table columns or to a new table

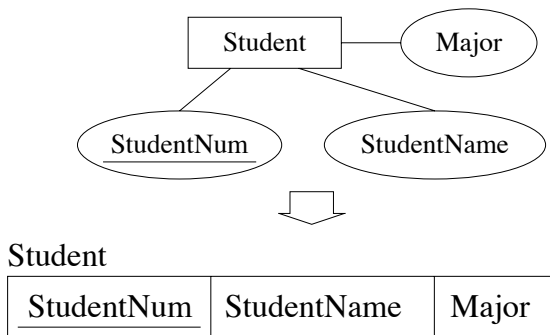
# Representing Strong Entity Sets

Entity set  $E$  with attributes  $a_1, \dots, a_n$  translates to table  $E$  with attributes  $a_1, \dots, a_n$

Entity of type  $E \leftrightarrow$  row in table  $E$

Primary key of entity set  $\rightarrow$  primary key of table

**Example:**



# Representing Weak Entity Sets

Weak entity set  $E$  translates to table  $E$

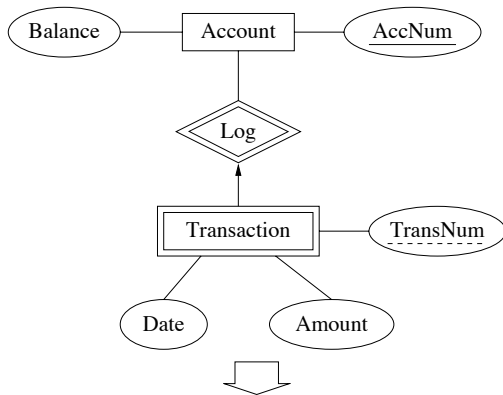
Columns of table  $E$  should include

- Attributes of the weak entity set
- Attributes of the identifying relationship set
- Primary key attributes of entity set for dominating entities

Primary key of weak entity set  $\rightarrow$  primary key of table

# Representing Weak Entity Sets (cont.)

## Example:



Account

<u>AccNum</u>	Balance
---------------	---------

Transaction

<u>TransNum</u>	<u>AccNum</u>	Date	Amount
-----------------	---------------	------	--------

# Representing Relationship Sets

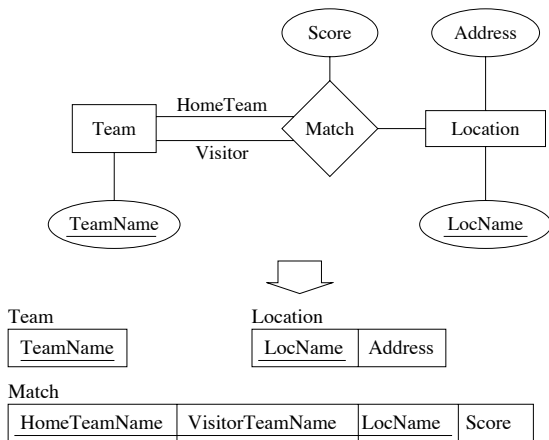
- If the relationship set is an identifying relationship set for a weak entity set then no action needed
- If we can deduce the general cardinality constraint (1,1) for a component entity set  $E$  then add following columns to table  $E$ 
  - Attributes of the relationship set
  - Primary key attributes of remaining component entity sets
- Otherwise: relationship set  $R \rightarrow$  table  $R$

# Representing Relationship Sets (cont.)

- Columns of table  $R$  should include
  - Attributes of the relationship set
  - Primary key attributes of each component entity set
- Primary key of table  $R$  determined as follows
  - If we can deduce the general cardinality constraint (0,1) for a component entity set  $E$ , then take the primary key attributes for  $E$
  - Otherwise, choose primary key attributes of each component entity

# Representing Relationship Sets (cont.)

## Example:



Note that the role name of a component entity set should be prepended to its primary key attributes, if supplied.



# Representing Aggregation

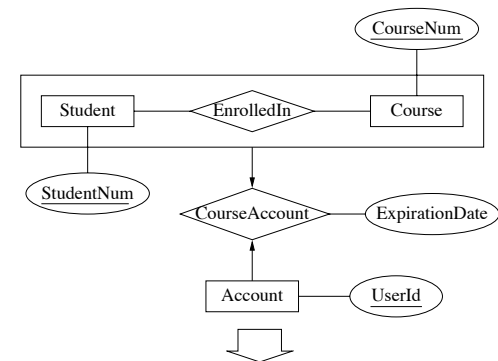
Tabular representation of aggregation of  $R$

= tabular representation for relationship set  $R$

To represent relationship set involving aggregation of  $R$ , treat the aggregation like an entity set whose primary key is the **primary key** of the table for  $R$

# Representing Aggregation (cont.)

## Example:



Student

<u>StudentNum</u>
-------------------

Course

<u>CourseNum</u>
------------------

Account

<u>UserId</u>
---------------

EnrolledIn

<u>StudentNum</u>	CourseNum
-------------------	-----------

CourseAccount

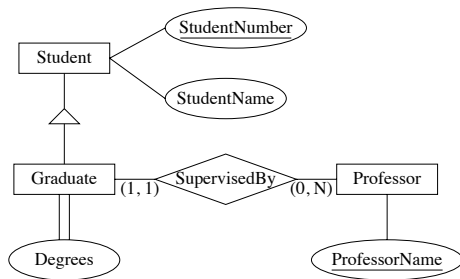
<u>UserId</u>	StudentNum	CourseNum	ExpirationDate
---------------	------------	-----------	----------------

primary key can be either UserId  
or the combination of StNum and CNum

# Representing Specialization

Create table for higher-level entity set, and treat specialized entity subsets like weak entity sets (without discriminators)

## Example:



Student

<u>StudentNumber</u>	StudentName
----------------------	-------------

Graduate

<u>StudentNumber</u>	ProfessorName
----------------------	---------------

Degree

<u>StudentNumber</u>	Degree
----------------------	--------

Professor

<u>ProfessorName</u>
----------------------

# Representing Generalization (Approach #1)

Create a table for each lower-level entity set only

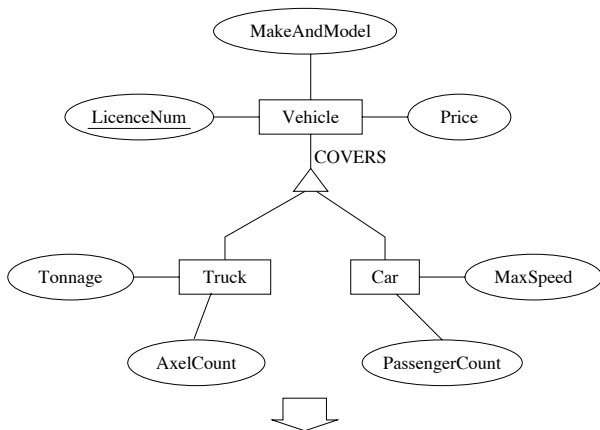
Columns of new tables should include

- Attributes of lower level entity set
- Attributes of the superset

The higher-level entity set can be defined as a view on the tables for the lower-level entity sets

# Representing Generalization (Approach #1)

## Example:



Truck

<u>LicenceNum</u>	MakeAndModel	Price	Tonnage	AxelCount
-------------------	--------------	-------	---------	-----------

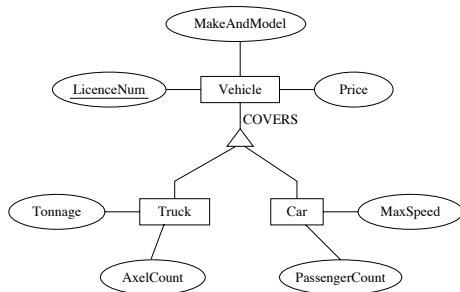
Car

<u>LicenceNum</u>	MakeAndModel	Price	MaxSpeed	PassengerCount
-------------------	--------------	-------	----------	----------------

# Representing Generalization (Approach #2)

Treat generalization the same as specialization.

## Example:



Vehicle

<u>LicenceNum</u>	MakeAndModel	Price
-------------------	--------------	-------

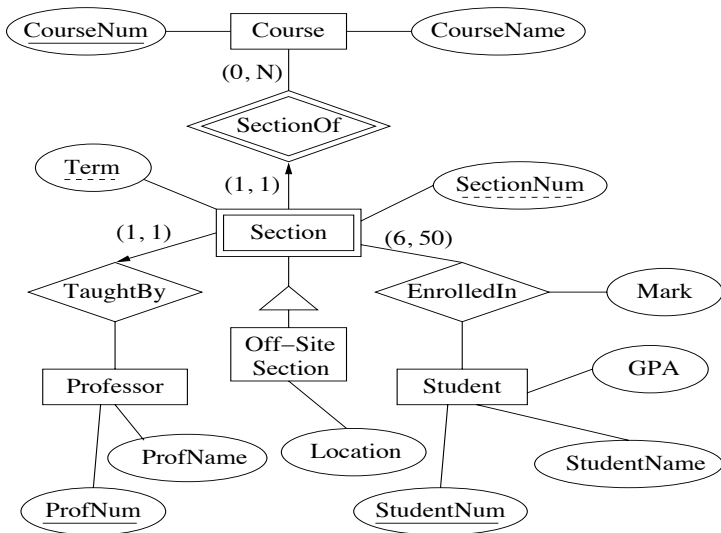
Truck

<u>LicenceNum</u>	Tonnage	AxelCount
-------------------	---------	-----------

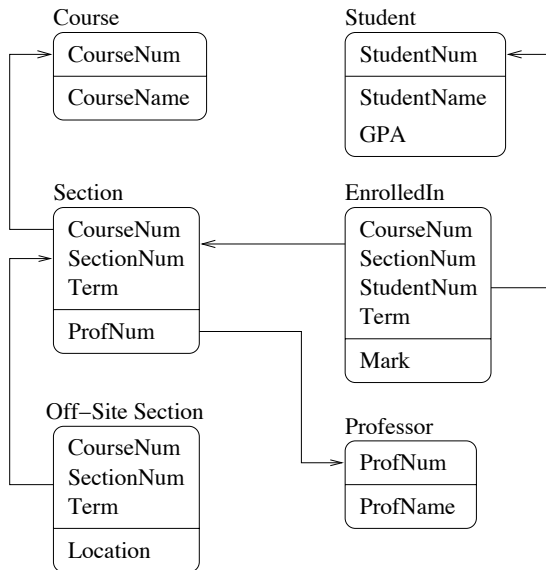
Car

<u>LicenceNum</u>	MaxSpeed	PassengerCount
-------------------	----------	----------------

# Example Translation: ER Diagram



# Example Translation: Relational Diagram





# Defining Relations and Integrity Constraints in SQL

- connected with a table (definition)

- ⇒ Primary Keys

- ⇒ Foreign Keys

- ⇒ CHECK constraints

- separate ECA rules (triggers)

# Tables with a Primary Key

- specifies a *attributes* of a table (w/types)
- specifies a *primary key* in a table
- syntax:

```
CREATE TABLE <name>
(
    ... <attributes>,
    PRIMARY KEY ( <list of attr> )
)
```

- also creates an unique index on the key

# Example

```
create table DEPT
( ID          integer not NULL,
  DeptName    char(20),
  MgrNO       char(3),
  PRIMARY KEY (ID)
)
```

## Example (cont.)

```
sql => insert into DEPT values \  
sql (cont.) => ( 1 , 'Computer Science', 000100)  
DB20000I The SQL command completed successfully.
```

```
sql => insert into DEPT values \  
sql (cont.) => ( 1 , 'Computer Science', 000100)  
SQL0803N One or more values in the INSERT or UPDATE  
statement are not valid because they would produce  
duplicate rows for a table with a unique index.  
SQLSTATE=23505
```

# Foreign Key

- specifies an *referential constraint*

- syntax:

```
CREATE TABLE <name>
(
    ... <attributes>,
    FOREIGN KEY ( <attrs> )
        REFERENCES <ref-table>( <attrs> )
        ON DELETE <delete-action>
        ON UPDATE <update action>
```

- the actions can be:

- \* RESTRICT – produce an error
- \* CASCADE – propagate the delete
- \* SET NULL – set to “unknown”

## Example

```
create table EMP
( SSN      integer not NULL,
  Name     char(20),
  Dept     integer,
  Salary   dec(8,2),
  primary key (SSN),
  foreign key (Dept) references DEPT(ID)
                        on delete cascade
                        on update restrict)
```

## Example (cont.)

```
db2 => insert into EMP \  
sql (cont.) => values ( 999, 'DAVE', 2, 50000 )  
SQL0530N  The insert or update value of FOREIGN KEY  
"DAVID.EMP.SQL970916001756640" is not equal to any  
value of the primary key of the parent table.  
SQLSTATE=23503
```

```
db2 => insert into EMP \  
sql (cont.) => values ( 999, 'DAVE', 1, 50000 )  
DB20000I  The SQL command completed successfully.  
db2 => select * from emp where SSN=999
```

SSN	NAME	DEPT	SALARY
999	DAVE	1	50000.00

## Example (cont.)

```
db2 => delete from DEPT where id=1
```

```
DB20000I  The SQL command completed successfully.
```

```
db2 => select * from emp where SSN=999
```

SSN	NAME	DEPT	SALARY
-----	-----	-----	-----



# CHECK constraints

- allow checking for “correct” data:
- syntax:

```
CREATE TABLE <name>
(
    ... <attributes>,
    CHECK <condition>
)
```

- condition is a *simple* search condition  
⇒ no subqueries (in DB2)

## Example

```
create table EMP
( SSN      integer not NULL,
  Name     char(20),
  Dept     integer,
  Salary   dec(8,2),
  primary key (SSN),
  foreign key (Dept) references DEPT(ID)
           on delete cascade
           on update restrict,
  check ( salary > 0 )
)
```

```
db2 => insert into emp values (998, 'DAVE', 1, 0 )
SQL0545N  The requested operation is not allowed
because a row does not satisfy the check constraint
"DAVID.EMP.SQL970916000939620".  SQLSTATE=23513
```

# Views

## Definition (View)

A *view* is a relation whose instance is determined by the instances of other relations.

A view has many of the same properties as a table.

- its schema information appears in the database schema
- access controls can be applied to it
- other views can be defined in terms of it

# Types of Views

- **Virtual:** Views are used only for querying; they are not stored in the database
- **Materialized:** The query that makes up the view is executed, the view constructed and stored in the database.

# SQL DDL: Views

## ■ General form:

```
CREATE VIEW <name>  
    [AS] ( <query> )
```

## ■ Example

```
create view ManufacturingProjects  
    ( Select projno, projname, firstnme, lastname  
      From project, employee  
      Where respemp = empno and deptno = 'D21' )
```

# Accessing a View

Query a view as if it were a base relation.

```
SELECT projname  
FROM    manufacturingprojects
```

What happens when you query a **virtual** view?

- At compile time, the view definition is found
- The query over the view is modified with the query definition
- The resulting query is optimized and executed

# Updating Views

- Modifications to a view's instance must be propagated back to instances of relations in conceptual schema.
- Some views cannot be updated unambiguously.

## Conceptual Schema

Persons

NAME	CITIZENSHIP
Ed	Canadian
Dave	Canadian
Wes	American

NationalPastimes

CITIZENSHIP	PASTIME
Canadian	Hockey
Canadian	Curling
American	Hockey
American	Baseball



## External Schema

PersonalPastimes

NAME	PASTIME
Ed	Hockey
Ed	Curling
Dave	Hockey
Dave	Curling
Wes	Hockey
Wes	Baseball

- 1 What does it mean to insert (Darryl, Hockey)?
- 2 What does it mean to delete (Dave, Curling)?

# View Updates in SQL

According to SQL-92, a view is updatable only if its definition satisfies a variety of conditions:

- The query references exactly one table
- The query only outputs simple attributes (no expressions)
- There is no grouping/aggregation/`DISTINCT`
- There are no nested queries
- There are no set operations

These rules are more restrictive than necessary.



# Materialized Views

## Problem

When a base table changes, the materialized view may also change.

- Solution?
  - Periodically reconstruct the materialized view.
  - Incrementally update the materialized view.
- Example: Data warehouses

# Data Control Language

assigns *access rights* to database objects

## ■ Syntax:

```
GRANT  <what> ON <object> TO <user(s)>  
REVOKE <what> ON <object> FROM <user(s)>
```

## ■ <what> ON <object> **can be**

```
DATABASE: BINDADD, CONNECT, CREATETAB  
          CREATE_NOT_FENCED, DBADM
```

```
INDEX: CONTROL
```

```
PACKAGE: BIND, CONTROL, EXECUTE
```

```
TABLE/VIEW: ALTER, CONTROL, INDEX, REFERENCES  
            SELECT, INSERT, DELETE, UPDATE
```

## ■ <user(s)> **is a list of**

(i) USER <name>      (ii) GROUP <name>      (iii) PUBLIC

# Summary

Schema design summary:

- 1 Create an ER diagram
  - ⇒ visualization of the design goals
- 2 Translate ER-to-Relational
- 3 Determine FD, MVD, JD, ...
  - ⇒ detect anomalies and decompose
  - ⇒ find **keys**
- 4 Determine inter-relational constraints
  - ⇒ INDs and foreign key constraints
- 5 Enforce rest of constraints
  - ⇒ `CHECK` declarations
  - ⇒ ECA rules (only as the last resort!)