

# Database Tuning and Physical Design: Execution of Transactions

Spring 2018

School of Computer Science  
University of Waterloo

Databases CS348

# Basics of Transaction Processing

Query (and update) processing converts requests for *sets of tuples* to requests for reads and writes of physical objects in the database.

database objects (depending on granularity) can be

- individual attributes
- records
- physical pages
- files (only for concurrency control purposes)

## Goals

- ⇒ correct and concurrent execution of queries and updates
- ⇒ guarantee that acknowledged updates are persistent

# ACID Requirements

Transactions are said to have the **ACID** properties:

**Atomicity:** all-or-nothing execution

**Consistency:** execution preserves database integrity

**Isolation:** transactions execute independently (as if they were executed in the system alone)

**Durability:** updates made by a committed transaction will not be destroyed by subsequent failures.

Implementation of transactions in a DBMS comes in two parts:

- **Concurrency Control:** committed transactions do not interfere
- **Recovery Management:** committed transactions are durable, aborted transactions have no effect on the database

# Concurrency Control: assumptions

- 1 we fix a database: a set of objects read/written by transactions:

$\Rightarrow r_i[x]$ : transaction  $T_i$  reads object  $x$

$\Rightarrow w_i[x]$ : transaction  $T_i$  writes (modifies) object  $x$

- 2 a transaction  $T_i$  is a sequence of operations

$$T_i = r_i[x_1], r_i[x_2], w_i[x_1], \dots, r_i[x_4], w_i[x_2], c_i$$

$c_i$  is the **commit request** of  $T_i$ .

- 3 for a **set of transactions**  $T_1, \dots, T_k$  we want to produce a *schedule*  $S$  of operations such that

$\Rightarrow$  every operation  $o_i \in T_i$  appears also in  $S$

$\Rightarrow T_i$ 's operations in  $S$  are ordered the same way as in  $T_i$

Goal:

produce a *correct schedule* with *maximal parallelism*

# Transactions and Schedules

If  $T_i$  and  $T_j$  are concurrent transactions, then it is always correct to schedule the operations in such a way that:

- $T_i$  will appear to precede  $T_j$  meaning that  $T_j$  will “see” all updates made by  $T_i$ , and  $T_i$  will not see any updates made by  $T_j$ , or
- $T_i$  will appear to follow  $T_j$ , meaning that  $T_i$  will see  $T_j$ 's updates and  $T_j$  will not see  $T_i$ 's.

Idea how to define Correctness:

it must appear as if the transactions have been executed sequentially (in some *serial* order).

# Serializable Schedules

## Definition

An execution of is said to be **serializable** if it is equivalent to a serial execution of the same transactions.

## Example:

- An interleaved execution of two transactions:

$$S_a = w_1[x] \ r_2[x] \ w_1[y] \ r_2[y]$$

- An equivalent serial execution ( $T_1$  ,  $T_2$ ):

$$S_b = w_1[x] \ w_1[y] \ r_2[x] \ r_2[y]$$

- An interleaved execution with no equivalent serial execution:

$$S_c = w_1[x] \ r_2[x] \ r_2[y] \ w_1[y]$$

# Conflict Equivalence

How do we determine if two schedules are *equivalent*?

⇒ cannot be based on any particular database instance

## Conflict Equivalence:

- two operations *conflict* if they
  - (1) belong to different transactions
  - (2) access the same data item  $x$
  - (3) at least one of them is a write operation  $w[x]$ .
- we require that in two *conflict-equivalent histories* all *conflicting operations* are ordered the same way.
- yields *conflict-serializable* schedules  
⇒ *conflict-equivalent* to a serial schedule

## View Equivalence:

allows more schedules, but it is harder (NP-hard) to compute

# Other Properties of Schedules

Serializability guarantees correctness. However, we'd like to avoid other **unpleasant** situations.

## Recoverable Schedules: (RC)

transaction  $T_j$  *reads* a value  $T_i$  has written,  $T_j$  succeeds to **commit**, and  $T_i$  tries to abort (in this order)

⇒ to abort  $T_2$  we need to *undo* effects of  
a *committed* transaction  $T_1$ .

⇒ commits only in order of the read-from dependency

## Cascadeless Schedules (ACA):

if  $T_j$  above didn't commit we can abort it:  
may lead to *cascading aborts* of many transactions

⇒ no reading of uncommitted data



# How to Get a Serializable Schedule?

So how do we build schedulers that produce serializable and cascadeless schedules?

The **scheduler** receives requests from the query processor(s). For each operation it chooses one of the following actions:

- execute it (by sending to a lower module),
- delay it (by inserting in some queue), or
- reject it (thereby causing abort of the transaction)
- ignore it (as it has no effect)

Two main kinds of schedulers:

- ⇒ conservative (favors delaying operations)
- ⇒ aggressive (favors rejecting operations)

# Two Phase Locking (2PL)

Transactions must have a **lock** on objects before access:

- a **shared lock** is required to read an object
- an **exclusive lock** is required to write an object

It is *insufficient* just to acquire a lock, access the data item, and then release it immediately...

## 2PL Protocol

A transaction has to **acquire** all locks before it **releases** any of them.

## Theorem

*Two-phase locking guarantees that the produced transaction schedules are (conflict) serializable.*

In practice: **STRICT 2PL** (locks held till commit; this guarantees ACA)

# Deadlocks and What to do

With 2PL we may end with a **deadlock**:

$r_1[x], r_2[y], w_2[x]$  (*blocked by  $T_1$* ),  $w_1[y]$  (*blocked by  $T_2$* )

How do we deal with this:

- deadlock prevention:

- ⇒ locks granted only if they can't lead to a deadlock.
- ⇒ ordered data items and locks granted in this order.

- deadlock detection:

- ⇒ wait for graphs and cycle detection.
- ⇒ resolution: the system **aborts** one of  
the offending transactions (involuntary abort).

in practice: detection (or often just a timeout) and abort

# Variations on Locking

- Multi-granularity Locking

- ⇒ not all locked objects have the same size
- ⇒ advantageous in presence of bulk vs. tiny updates

- Predicate Locking

- ⇒ locks based on selection predicate rather than on a value

- Tree Locking

- ⇒ tries to avoid congestion in roots of (B-)trees
- ⇒ allows relaxation of 2PL due to tree structure of data

- Lock Upgrade protocols

- ...

# Inserts and Deletes

We have been assuming a **fixed set** of data items.

⇒ what if we try to *insert* or *delete* an item?

■ does plain 2PL (correctly) handle this situation? NO:

⇒ one transaction tries to count records in a table

⇒ second transactions adds/ deletes a record

■ this situation is called the **phantom problem**.

Solution: operations that ask for “all records” have to lock  
against insertion/deletion of a qualifying record

⇒ locks on tables

⇒ index locking and other techniques

# Isolation Levels in SQL

The guarantee of serializable executions may carry a heavy price. Performance may be poor because of blocked transactions and deadlocks.

Four **isolation levels** are supported:

**Level 3:** (Serializability)

⇒ essentially table-level strict 2PL

**Level 2:** (Repeatable Read)

⇒ tuple-level strict 2PL; “phantom tuples” may occur

**Level 1:** (Cursor Stability)

⇒ tuple-level exclusive-lock only strict 2PL

reading the same object twice: different values

**Level 0:**

⇒ neither read nor write locks are acquired

⇒ transaction may read uncommitted updates

# Recovery: Goals and Setting

Two goals:

- 1 allow transactions to be  
**committed** (with a guarantee that the effects are permanent) or  
**aborted** (with a guarantee that the effects disappear)
- 2 allow the database to be **recovered** to a consistent state in case on HW/power/. . . failure.

**Input:** a *2PL*, *ACA* schedule of operations produced by TM.

**Output:** a schedule of reads/writes/**forced writes**.

# Approaches to Recovery

Two essential approaches:

## 1 Shadowing

- ⇒ copy-on-write and merge-on-commit approach
- ⇒ poor clustering
- ⇒ used in system R, but not in modern systems

## 2 Logging

- ⇒ use of LOG (separate disk) to avoid forced writes
- ⇒ good utilization of buffers
- ⇒ preserves original clusters



# Log-Based Approaches

A log is a read/**append only** data structure (a file)

⇒ transactions add **log records** about what they do

Log records contain several types of information:

- **UNDO information:** old versions of objects that have been modified by a transaction. UNDO information can be used to undo database changes made by a transaction that aborts.
- **REDO information:** new versions of objects that have been modified by a transaction. REDO records can be used to redo the work done by a transaction that commits.
- **BEGIN/COMMIT/ABORT** records are recorded whenever a transaction begins, commits, or aborts.

# Example of a LOG

log head	→	$T_0, \text{begin}$
(oldest part)		$T_0, X, 99, 100$
		$T_1, \text{begin}$
		$T_1, Y, 199, 200$
		$T_2, \text{begin}$
		$T_2, Z, 51, 50$
		$T_1, M, 1000, 10$
		$T_1, \text{commit}$
		$T_3, \text{begin}$
		$T_2, \text{abort}$
		$T_3, Y, 200, 50$
		$T_4, \text{begin}$
(newest part)		$T_4, M, 10, 100$
log tail	→	$T_3, \text{commit}$

# Write-Ahead Logging (WAL)

How do we make sure the LOG is consistent with the main database?

Write-Ahead Logging (WAL) approach requires:

- 1 UNDO rule:** a **log record** for an update is written to log disk **before** the corresponding data (page) is written to the *main* disk  
(guarantees *Atomicity*)
- 2 REDO rule:** **all log records** for a transaction are written to log disk before **commit**  
(guarantees *Durability*)

# Summary

ACID properties of transactions guarantee correctness of concurrent access to the database and of data storage.

- consistency and isolation based on **serializability**
  - ⇒ leads to definition of correct **schedulers**
  - ⇒ responsibility of the **transaction manager**
- durability and atomicity
  - ⇒ responsibility of the **recovery manager**
  - ⇒ synchronous writing is too inefficient
    - replaced by synchronous writes to a LOG and WAL